

Advanced School on  
**QUANTUM MODELLING**  
of Materials with CRYSTAL



UNIVERSITÀ  
DI TORINO

# Parallel Computing in the CRYSTAL Code

---

Giacomo Ambrogio

Volta Redonda, State of Rio de Janeiro  
Brazil

[giacomo.ambrogio@unito.it](mailto:giacomo.ambrogio@unito.it)

QMMC 2026

# Outline

- **Why parallel?**
  - The serial problem
  - Parallel **hardware** architectures
- **When to use parallel computing**
  - The **overhead**
  - Parallel Implementations in Crystal
    - The algorithm (SCF)
    - **PCRYSTAL**
    - **MPPCRYSTAL**
    - Examples
    - **OpenMP**
  - Using GPUs: **GCRYSTAL**

# Why Parallel?

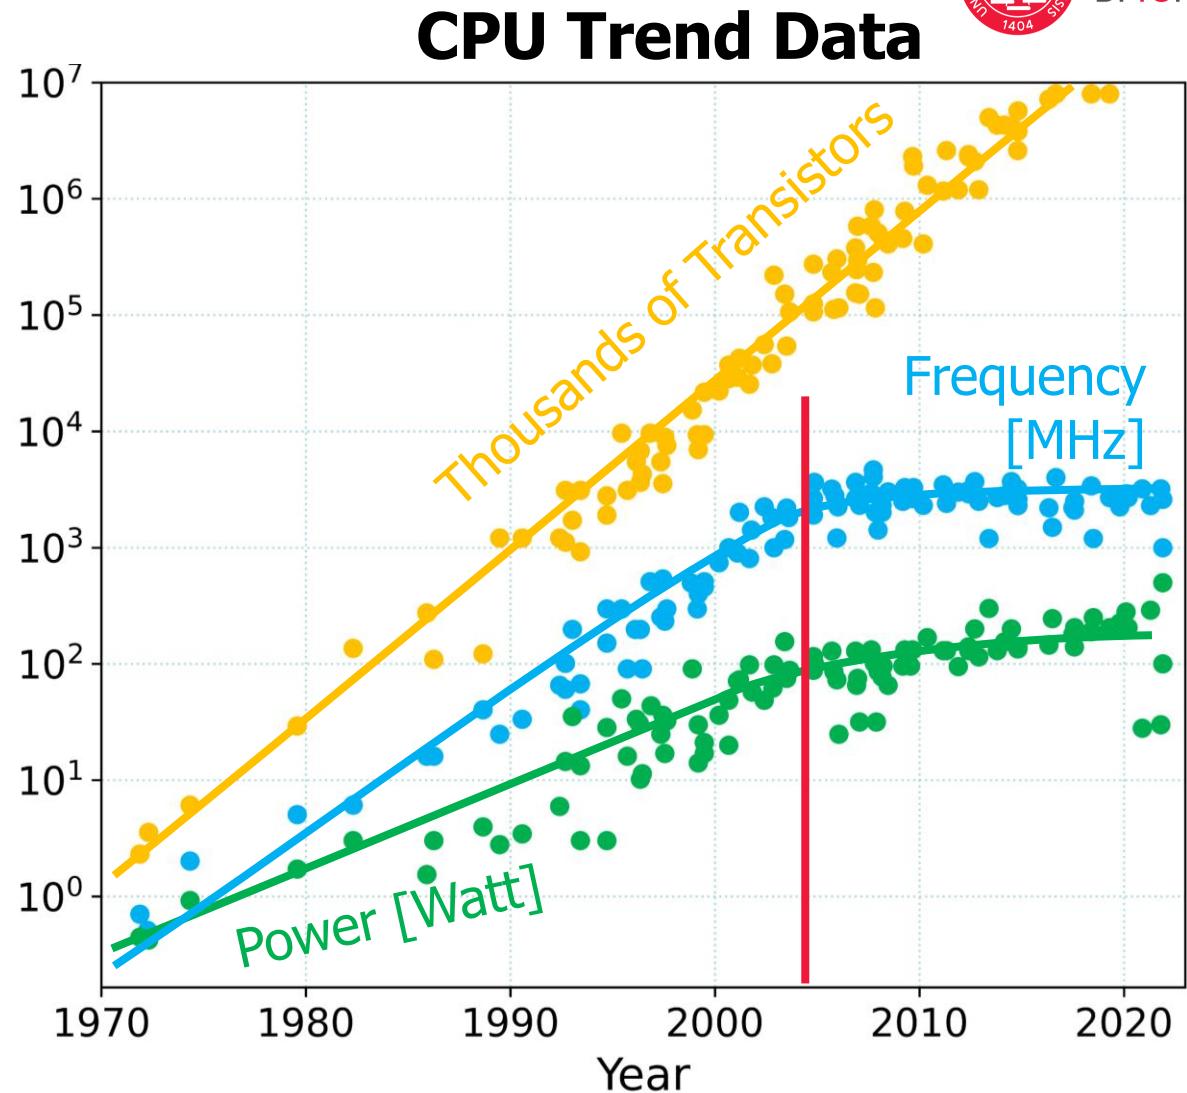
- Why do we need parallel computers?
- Why should you be interested in using parallel computers?

# The Serial Problem

**Moore's Law:** number of **transistors** in microchips *doubles* roughly every two years

At first, we were consistently able to increase the **frequency** (operations performed per second).

Around 2002, hardware reached a **physical limitation**  
increasing power → silicon chip would melt



# The Serial Problem

Some crazy attempts to keep low temperature: **Liquid Nitrogen**



Intel  
i9-12900K  
**6.9 GHz**  
2022



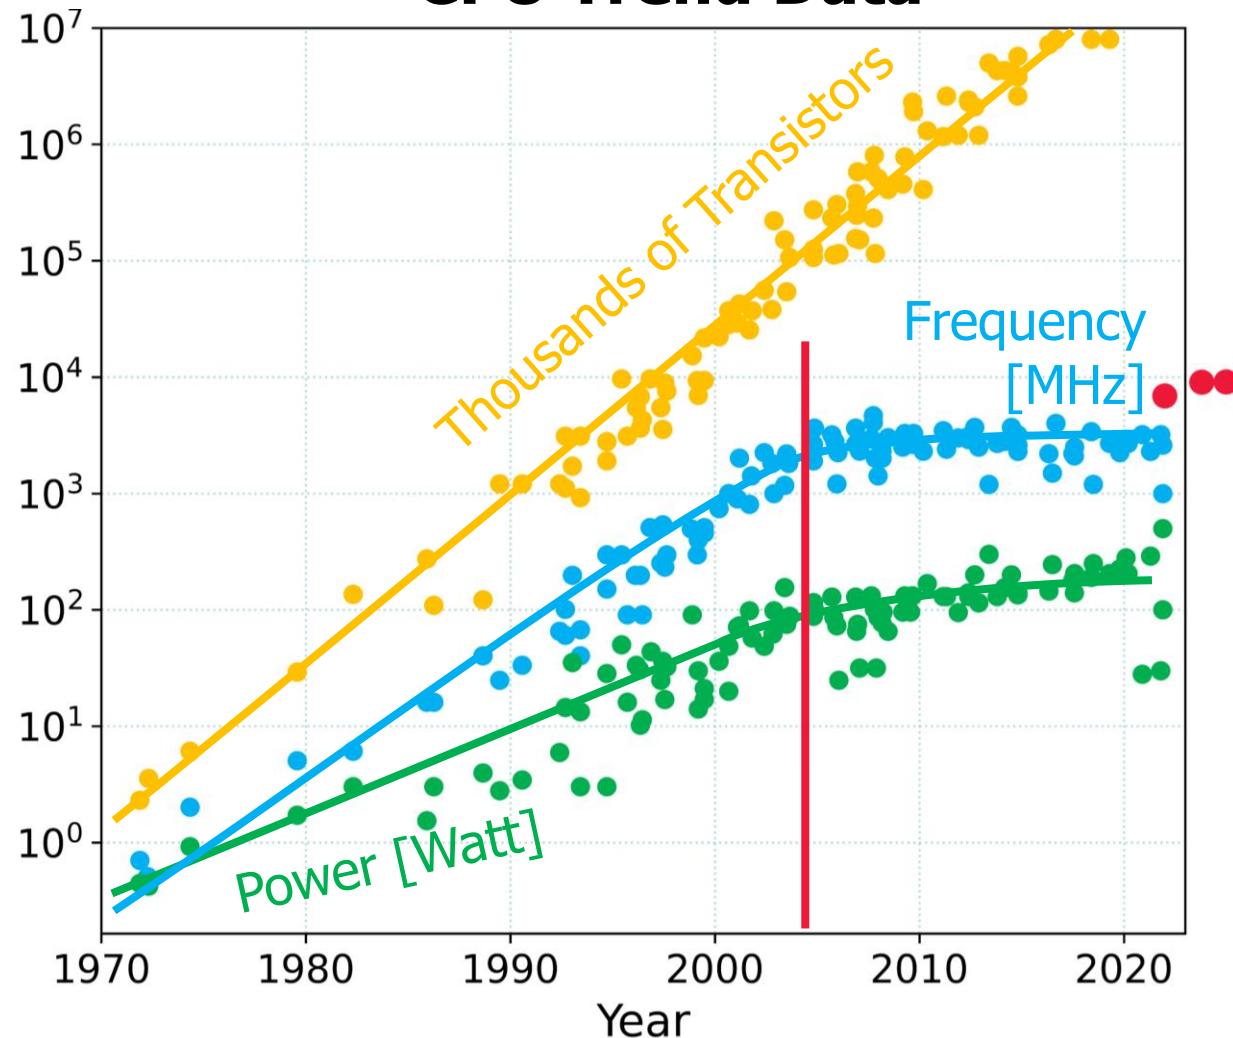
Intel  
i9-14900K  
**9.04 GHz**  
2023



Intel  
i9-14900KS  
**9.12 GHz**  
2025

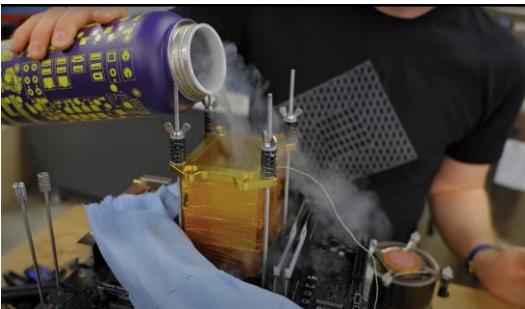


## CPU Trend Data



# The Serial Problem

Some crazy attempts to keep low temperature: **Liquid Nitrogen**



Intel  
i9-12900K  
**6.9 GHz**  
2022



Intel  
i9-14900K  
**9.04 GHz**  
2023

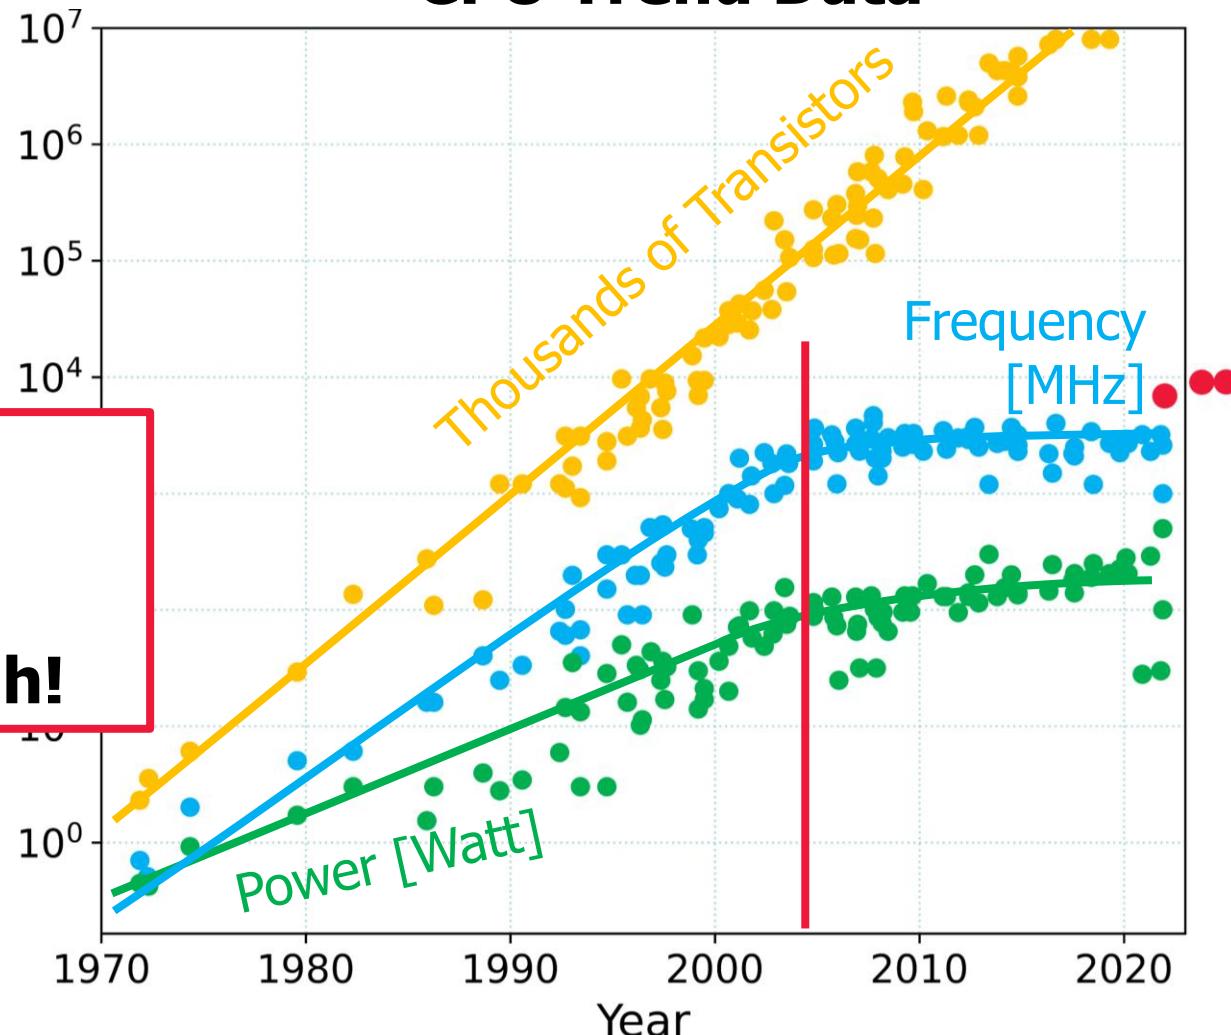


Intel  
i9-14900KS  
**9.12 GHz**  
2025



- Impractical
- Unstable
- **Not Enough!**

**CPU Trend Data**

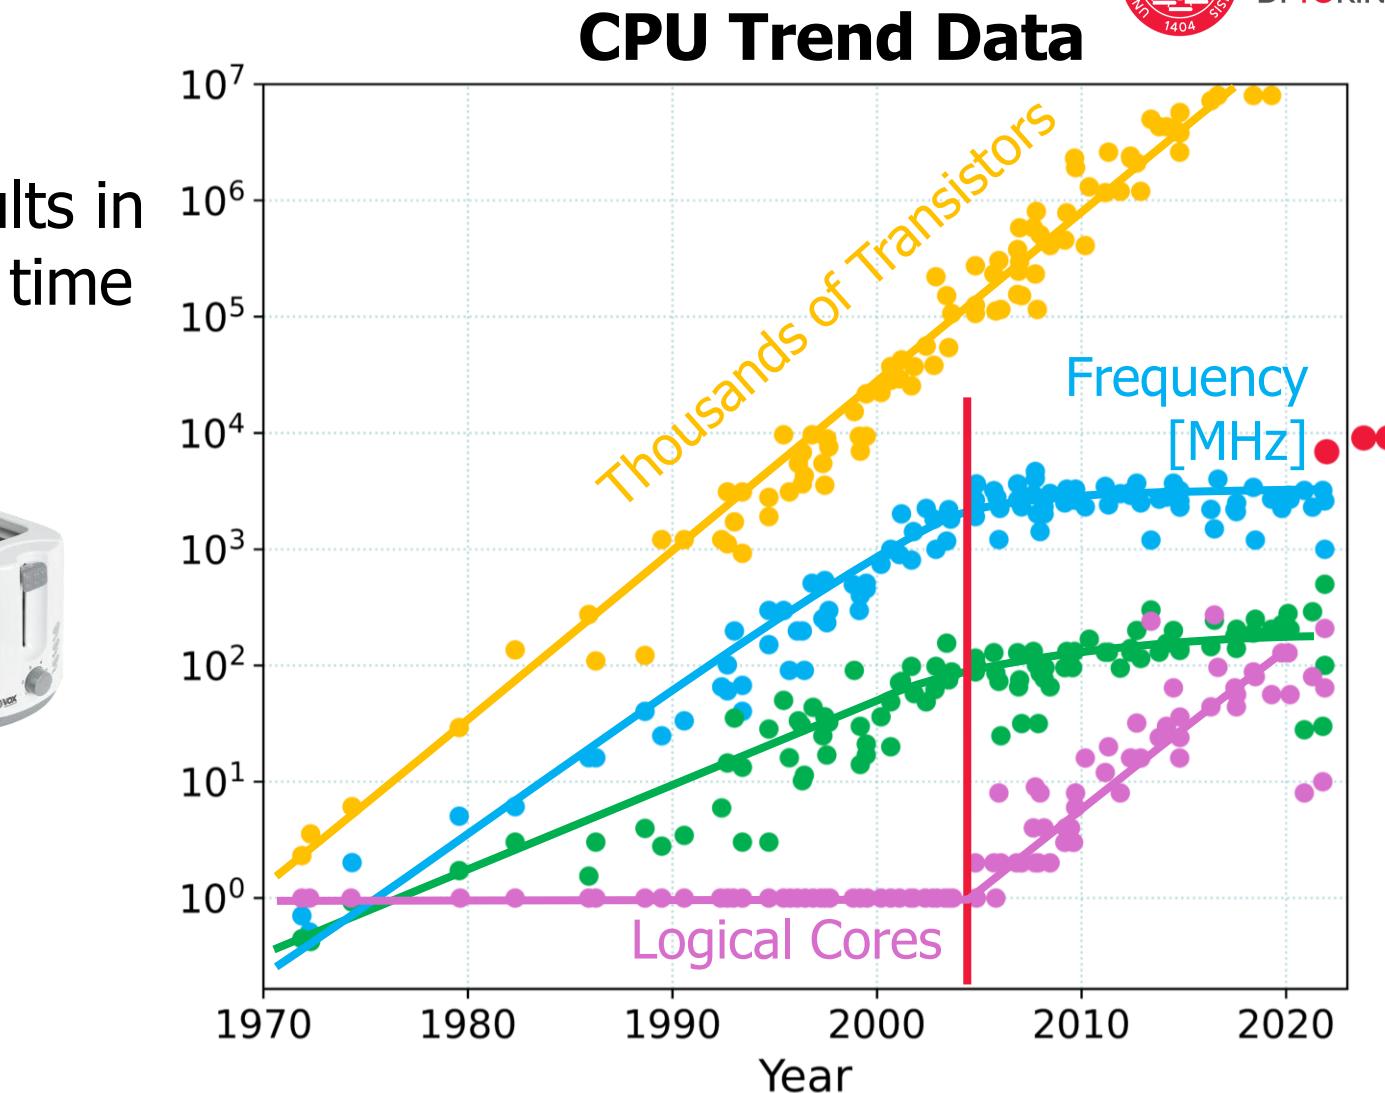


# The Serial Problem

The solution is to go **Parallel**

More computing units  
(in the same chip) → More results in  
**CORES** the same time

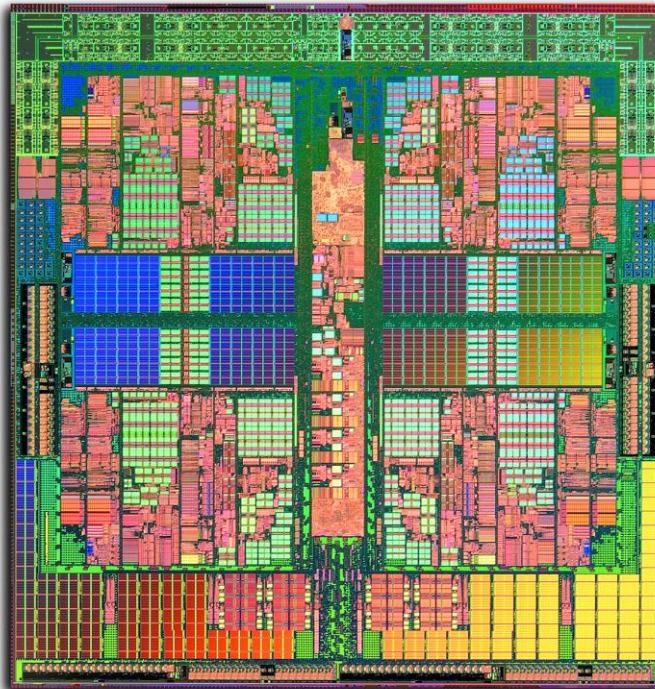
All modern computers and  
electronic devices are parallel:



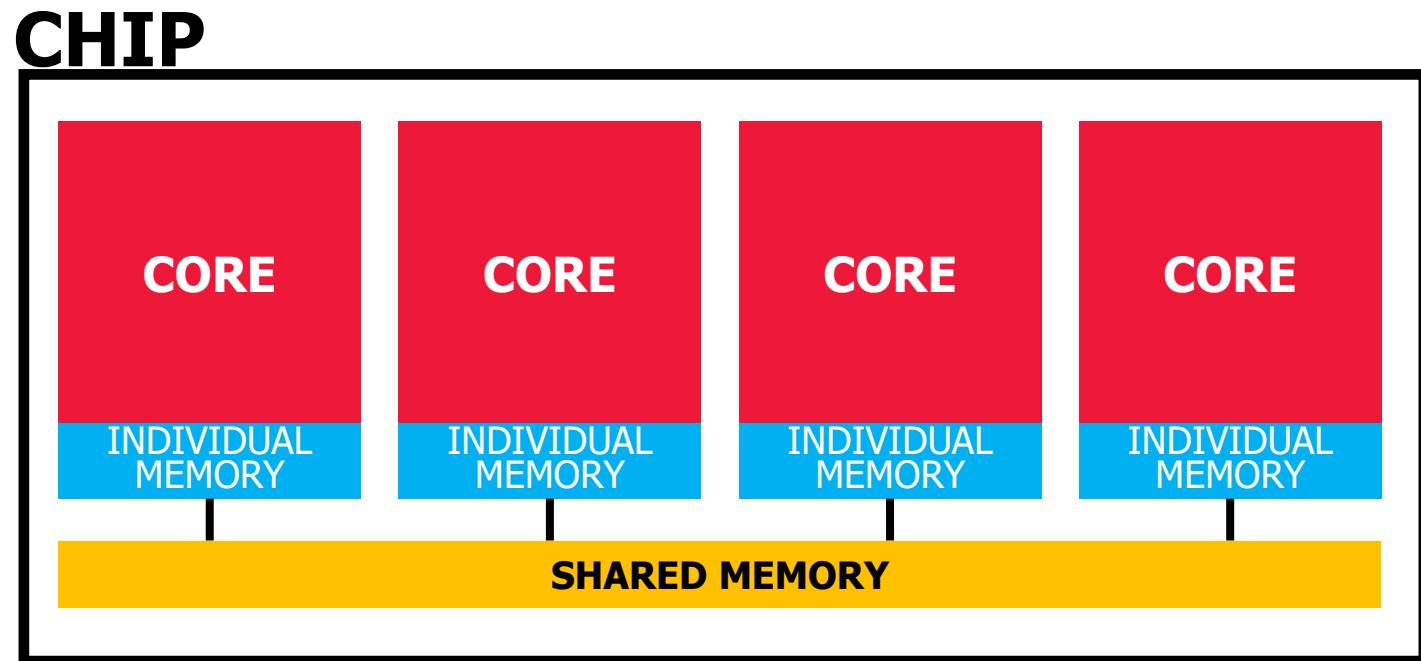
# Inside Parallel Computers

The main component is the  
**Multicore CPU**

→ From 2 to **192** cores per Chip!



Quad-Core AMD  
Opteron  
2007



Off-chip components

# Inside Large Parallel Computers

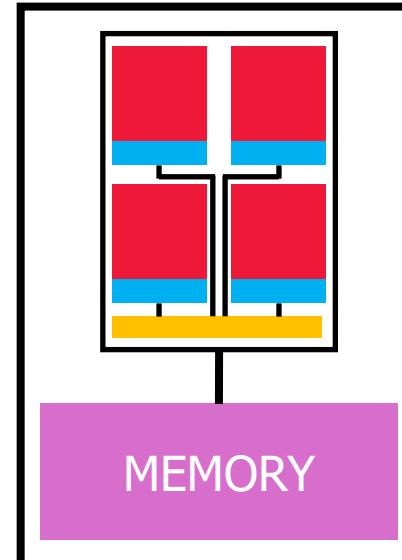
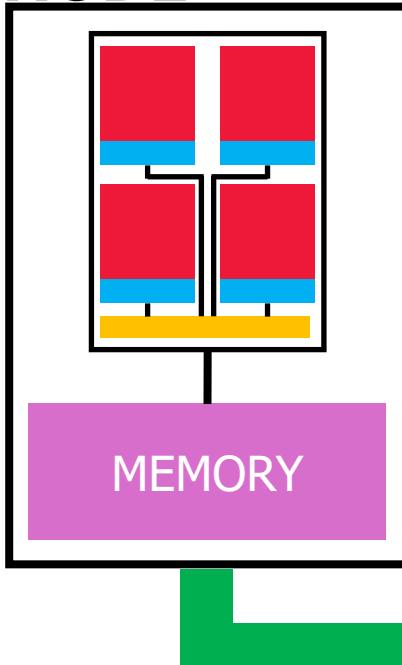
Multicore processors are clustered together with an **interconnect**

→ ethernet, InfiniBand

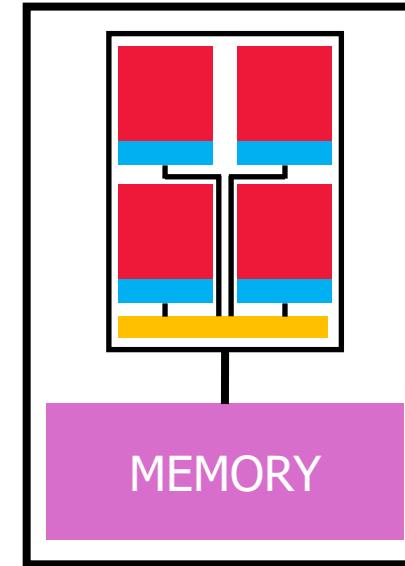
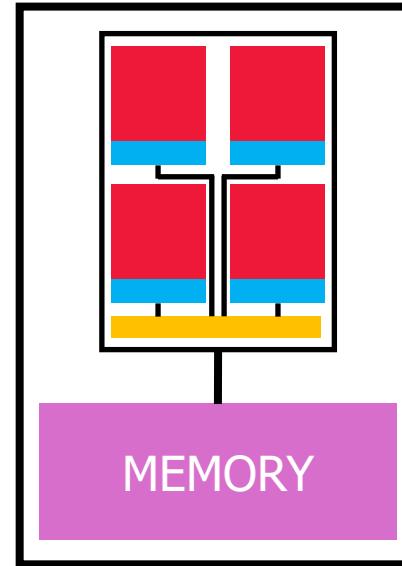
**Supercomputer**  
High Performance Computing



**NODE**



...



Same components as in small systems

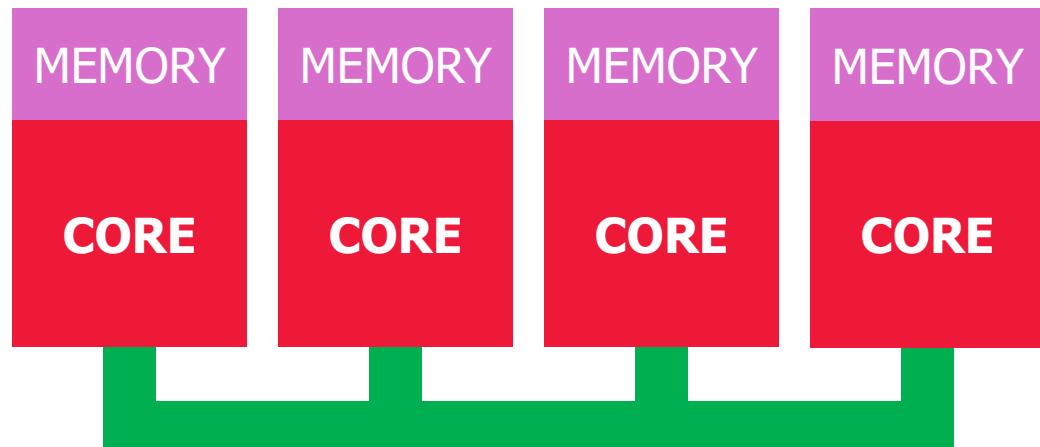
- More cores
- More memory

Very expensive!

# Simplify the Model (for now)

Consider a simplified view: a number of **cores** connected via an **interconnect**, each with its own **private memory**

- Ignore multicore
- Ignore limits of interconnection



# Why Parallel?

- Faster **time to solution**  
Jobs that takes days or weeks → hours
- More **Memory** so larger jobs  
32 nodes means 32 times more memory  
Hopefully, one can run a **larger jobs**
- **BUT** more complex **to run**
- Depends on **local setup**
- And more **difficult** to get the best out of the program



I will try to give some **tips** on how to do this!

# Always Parallel?

You *probably* always want to run parallel. However,

- **10 people** do not necessarily get a **job done 10 times quicker**  
→ You need **enough work** to keep 10 people busy!
- Similarly, **10 cores** does not mean the computation will run **10 times quicker**  
→ You need to have enough work to **keep 10 cores busy!**

**Do not always use all available cores**

- There might be sufficient work for 10 cores, but not for 100
- **Overhead:** using more cores might make your program SLOW DOWN

# The Overhead

## Work

$$\sim O\left(\frac{N}{P}\right)$$

## Overhead

$$\sim O(P)$$

$P$  Number of cores

Reasons for overhead in parallel programs include:

- **Amdahl's law:** not all the program is parallelised
- **Load balance:** not all tasks are equal
- The **interconnect** (necessary, but it is slow)
- Memory **bandwidth**
- **I/O**
- O/S jitter
- Phase of the moon
- Black magic
- The computer hates you

# The Overhead

## Work

$$\sim O\left(\frac{N}{P}\right)$$

As  $N$  increases the overhead becomes **less important**

## Overhead

$$\sim O(P)$$

$P$  Number of cores

**Parallel Computing Is Best For Solving Large Systems**

# Trying To Quantify It

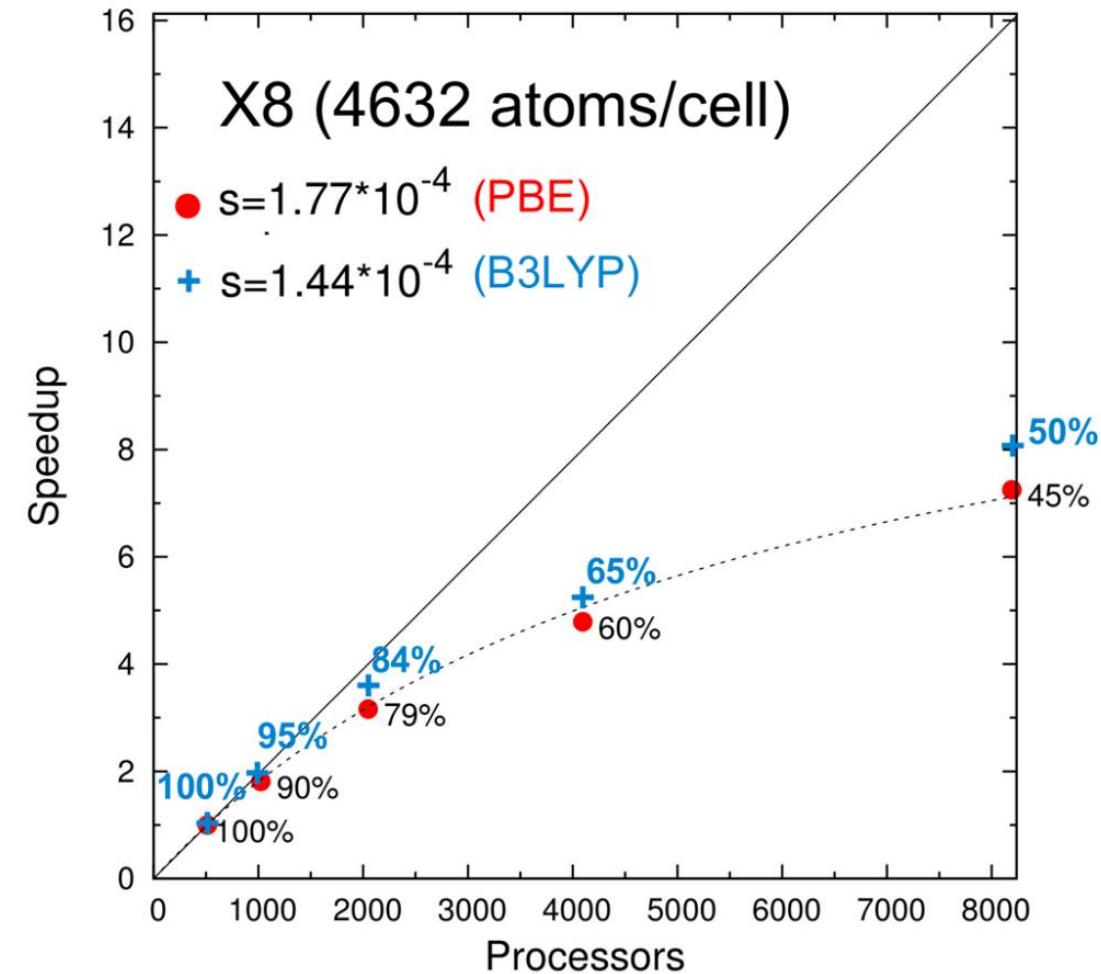
**Each job** have its own **characteristics**

→ may be best to run on a  
**different number of cores**

What do we mean by “best”?

**Need to measure something!**

- **Speed up:** applies to any code
- **CRYSTAL:** SCF cycles per unit time



# Parallel Implementations in CRYSTAL

## PCRYSTAL

- **Replicated data**
- All of CRYSTAL implemented
- Good for **small-medium systems** on **small-medium core counts**
- *Reference Implementation*

## MPPCRYSTAL

- **Distributed data**
- Most of CRYSTAL implemented
- Good for **medium-large problems** on **medium-very large core counts**

## OpenMP

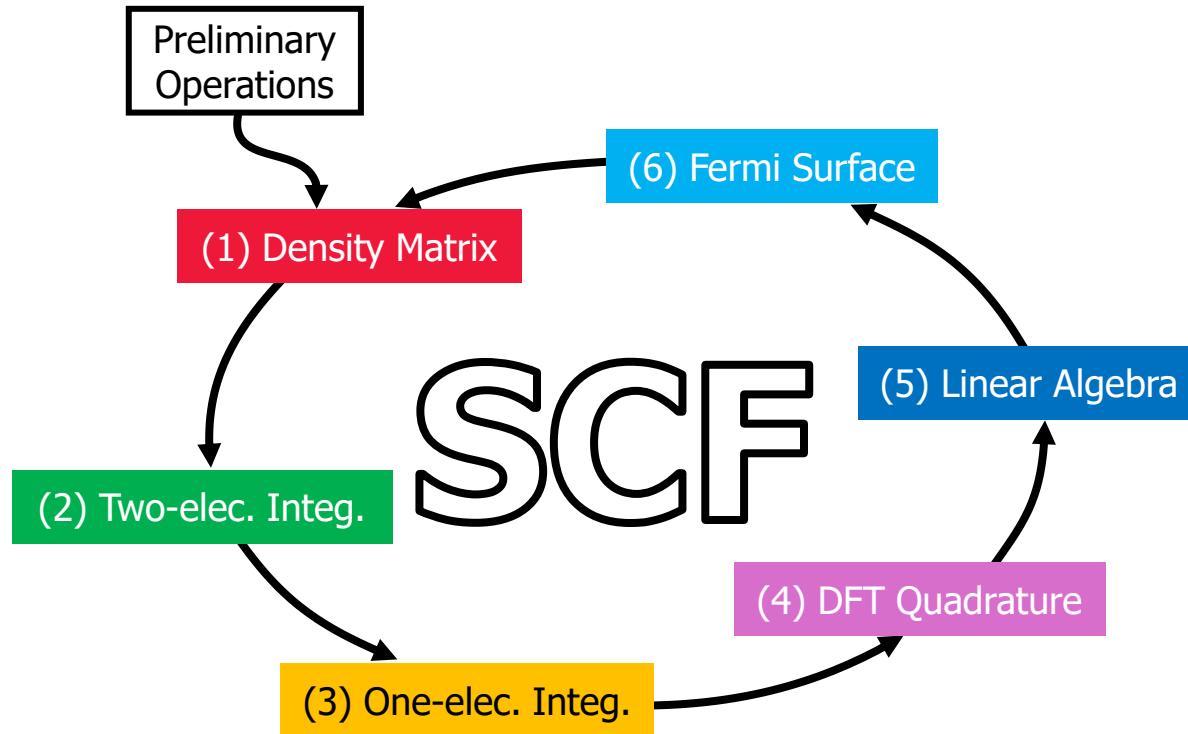
- New in CRYSTAL23
- More later

## GCRYSTAL

- In **development** for **next release**
- More later

# Basic Algorithm

Single point **energy**  
(and **wavefunction**) calculation:



Repeat until **convergence** on energy

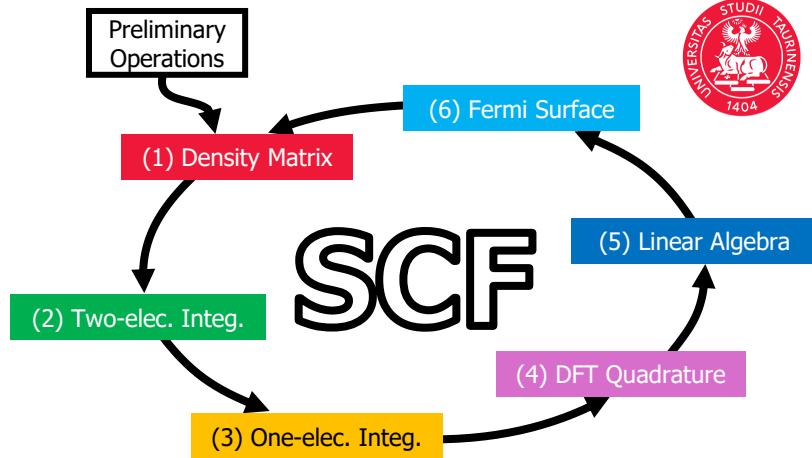
1. Start with guess of **Density matrix**
- 2./3. The **Hamiltonian** matrix  $H_g$  is obtained as a sum of **independent integrals**
4. **DFT terms** are added to  $H_g$
5.  $H_g$  is converted in reciprocal space:  
set of  $H_k$  matrices obtained.  
 $H_k$  is orthogonalized.  
**Solve for**  $H_k C_k = S_k C_k E_k$
6. Compute **Fermi Surface**
1. Form new **Density Matrix** from  $C_k$

Based on  MPI (4.x.x) standard for parallelism  
Should compile and work **everywhere**

## Replicated Data

- Each processor has a **complete copy** of “all” objects used by the code
- Make **implementation simple**
- But memory limits size of system that can be studied

Based on  MPI (4.x.x) standard for parallelism  
 Should compile and work **everywhere**



## Replicated Data

- Each processor has a **complete copy** of “all” objects used by the code
- Make **implementation simple**
- But memory limits size of system that can be studied

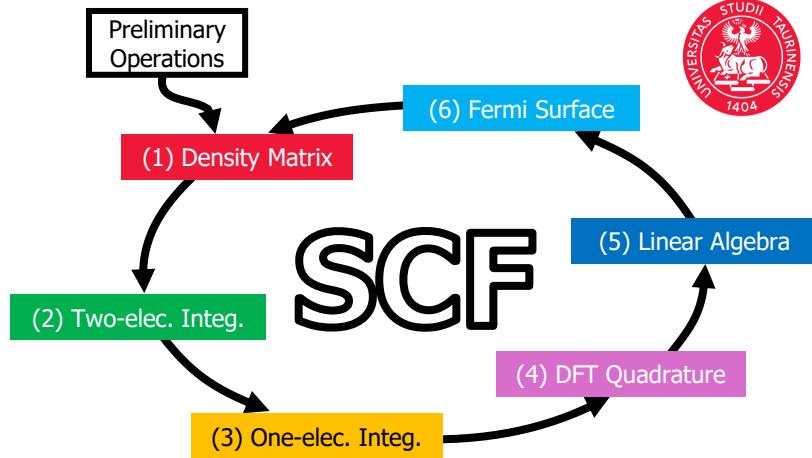
## Generating the Hamiltonian $H_g$ matrix

- Each integral is **independent**
  - give a subset to each one of the cores
- DFT: **integration** over a **grid** of independent points
  - each core computes a different set of points

**Almost perfectly parallel!**

→ Limit on scaling is load balancing

Based on  MPI (4.x.x) standard for parallelism  
 Should compile and work **everywhere**



## Replicated Data

- Each processor has a **complete copy** of “all” objects used by the code
- Make **implementation simple**
- But memory limits size of system that can be studied

Solving Secular Equation  $\mathbf{H}_k \mathbf{C}_k = \mathbf{S}_k \mathbf{C}_k \mathbf{E}_k$

- Each **k** point is **independent**  
 → each core can deal with a subset

**But number of **k** points limits parallelism**

Large systems require few **k** points (only 1 in the limit)

Each core **allocate one full matrix** in memory  
 → **Limit in size despite number of cores used**

# PCRYSTAL – How to run

Ultimately depends on your **local setup**

General command line is:

```
>> mpirun -np 4 /path/to/Pcrystal
```

Usually, you will need a **batch script** if running on a cluster

Strictly **NO changes** required to INPUT file

But it *must* be named INPUT, and be located in same directory where you launch the code

# PCRYSTAL – Summary

In general **scales well** where integrals dominate  
(small-medium systems)

**k** points **limit scaling** in linear algebra section

**Limit on size due to memory/core**  
(same as serial CRYSTAL)

→ Some tricks can be used to partially overcome this

# MPPCRYSTAL

Also based on  (4.x.x) standard for parallelism

Exploit **ScaLAPACK library** for linear algebra

[www.netlib.org/scalapack/](http://www.netlib.org/scalapack/) (requirement)

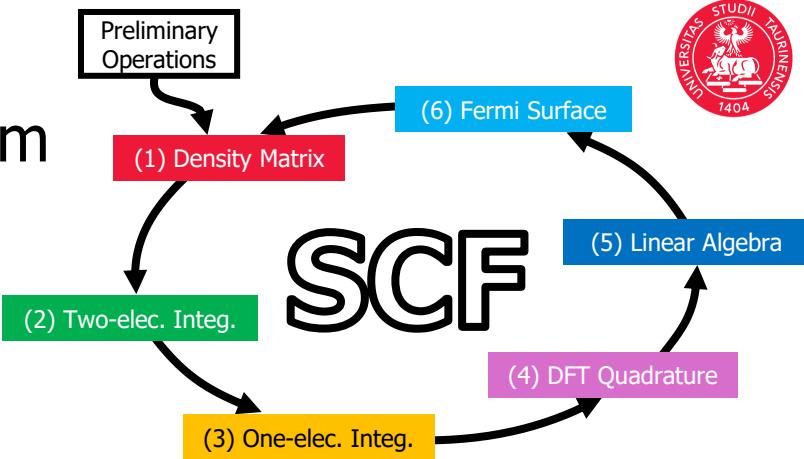
## Distributed Data

- Each core holds only a **part** of large data structures
- More **complex to implement**
- More cores and memory → **Larger jobs**

# MPPCRYSTAL

Also based on **MPI** (4.x.x) standard for parallelism

Exploit **ScalAPACK library** for linear algebra  
[www.netlib.org/scalapack/](http://www.netlib.org/scalapack/) (requirement)



## Distributed Data

- Each core holds only a **part** of large data structures
- More **complex to implement**
- More cores and memory → **Larger jobs**

Construction of  $H_g$  matrix same as PCRYSTAL

## Linear Algebra

- Exploit **k** point **parallelism**
- Each **k** point exploit **further levels of parallelism**
- More communication → **More overhead**  
Work scales  $O(N^3)$ , communications scale  $O(N^2)$
- **No I/O** (just store everything in memory)

# MPPCRYSTAL – Summary

Good for **larger systems** and **large core counts**

**Overhead** can dominate in smaller systems

→ **Rough estimate:** Use MPP if  $\frac{\text{spins} \times \text{k points} \times \text{AO}}{\text{cores}} \geq 20$   
(Very machine and case dependent)

**Distributed data** means that given enough cores can solve *any* system

**Not all of CRYSTAL is implemented**

(Will fail quickly and cleanly if the requested feature is not available)

Running is just the same as PCRYSTAL

# Properties

Both PCRYSTAL and MPPCRYSTAL write **final output** files (fort.9 i.e. wavefunction) that are in **exactly the same** form as the serial code

- **Run PROPERTIES just as before**
- But might take a long time and take a lot of memory

→ **PROPERTIES**, *replicated* data like PCRYSTAL

→ **MPPP\_PROPERTIES**, *Distributed* data like MPPCRYSTAL

- Supports NEWK, BAND, DOS, ECH3, POT3

# How To Get The Best Out Of Them

Basic procedure is:

- A series of **short runs** (small number of SCF cycles) on different number of cores to look at the scaling
- Look at the output to see if you can work around any **warnings\***
- **New set of short runs with improved input**

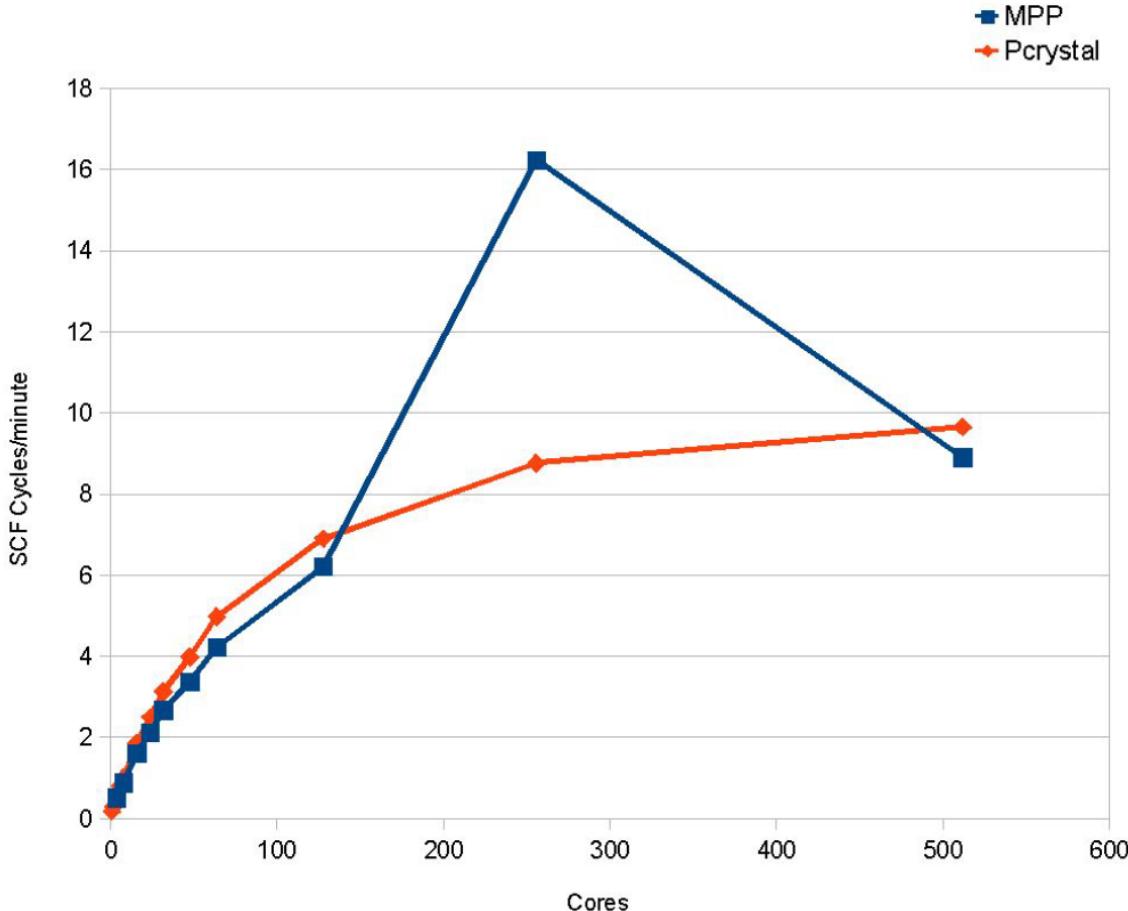
From that data **select PCRYSTAL or MPPCRYSTAL** and the number of cores to be used

\***Coulomb and Exchange buffers**

\***MPP allk/splitk** (see later)

# Very Old Example

- Organic crystal used in a drug study
- **624 basis functions**
- **8 irreducible k points**
- Calculation: **geometry optimisation**  
→ For performance tests just run the first SCF
- 16 cycles



# Very Old Example

So far **PCRYSTAL** looks better

→ System **too small** for diagonalisation  
in MPPCRYSTAL to work well

**EXCEPT at 256 cores.** What's going on?

At **128** cores we see in the output:

```
WARNING **** MPP **** ALL K POINTS DONE BY ALL PROCESSORS
```

**ALLK**

At **256** we see:

```
INFORMATION **** MPP **** MPP running in k point/spin parallel mode
```

**SPLITK**

→ **Extra performance** due to deciding a *good* number of cores  
for **k** point parallelism (SPLITK)

# The Art of Gentle Persuasion

How to force it to go k-point parallelism?

- In the last section put:

CMPLXFAC

2



This means you estimate your complex **k** points are **twice as expensive** as purely real ones

- Then look at the nature of your k-points

```
*** K POINTS COORDINATES (OBLIQUE COORDINATES IN UNITS OF IS = 3)
1-R( 0 0 0) 2-C( 1 0 0) 3-C( 0 1 0) 4-C( 1 1 0)
5-C( 0 0 1) 6-C( 1 0 1) 7-C( 0 1 1) 8-C( 1 1 1)
```

we have **1 real k point** and **7 complex**: the **unit cost** is  $2 \times 7 + 1 = 15$

- So run on **multiples** of **15 cores** and the code should decide to go **SPLITK**

# A Bit of Help

If running in **ALLK**, the code print out a **message** similar to the below to help you (*n.b. Not for the example discussed above*):

\*\*\*\*\*  
\*\*\*\*\*  
\*\*\*\*\*

Please, consider that for this calculation a double level of parallelism can be exploited by using a number of processors given by:

`N_proc = 44 x a_non_prime_integer (as 4, or 6, or 8, or 9, ...)`

For instance, `N_proc = 176, or 264, or 352, or 396, or ...`

`N_proc = 2816, or 5632, or 11264, or 22528, or ...`

`N_proc = 45056, or 90112, or ...`

\*\*\*\*\*  
\*\*\*\*\*  
\*\*\*\*\*

# New Performance

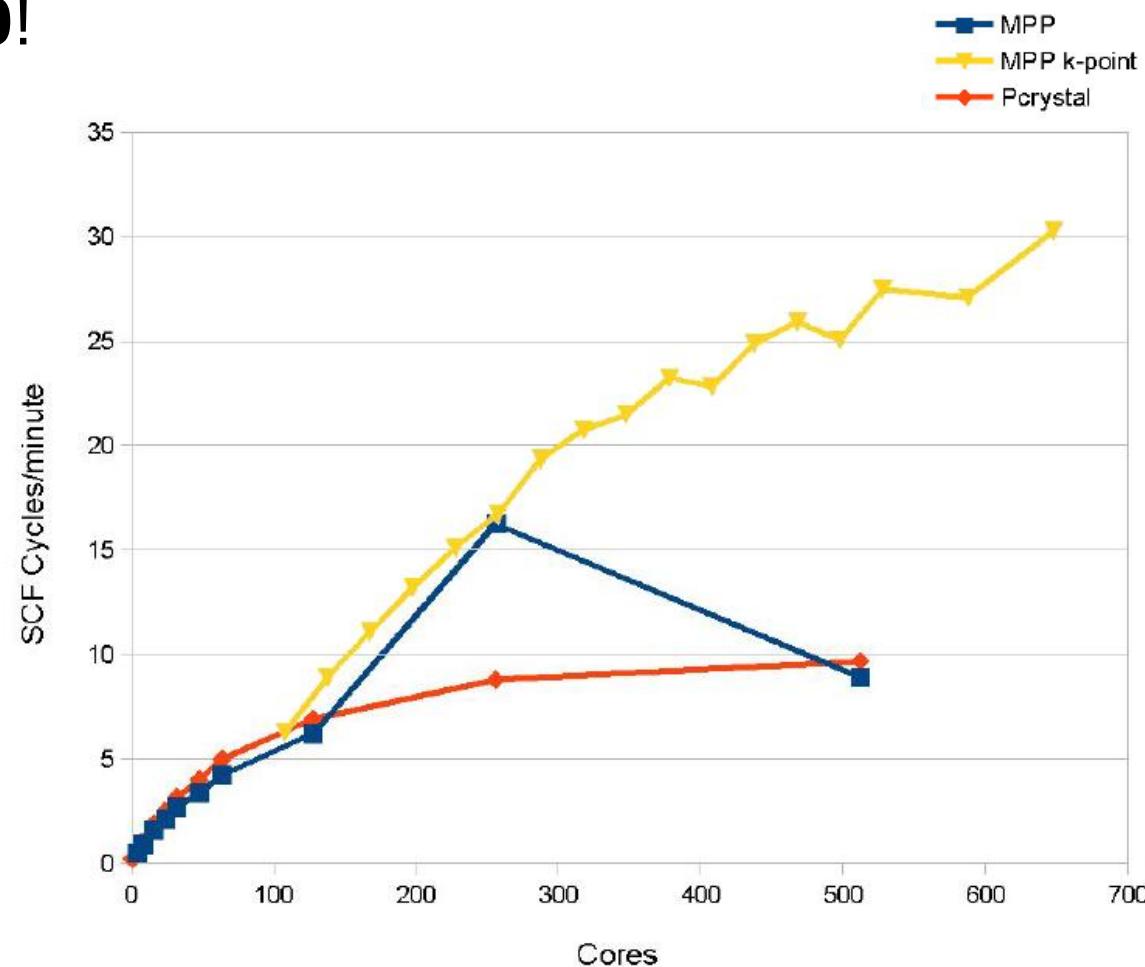
With a bit of work, we have increased our number of SCF cycles from 10 per minute to **around 30!**

- And the best code is now MPPcrystal

Of course **scaling is not perfect**, up to you to make decision where is the best for you, especially if **paying for the time**

Rough estimate above gives 250 cores for MPP

- Not too bad!

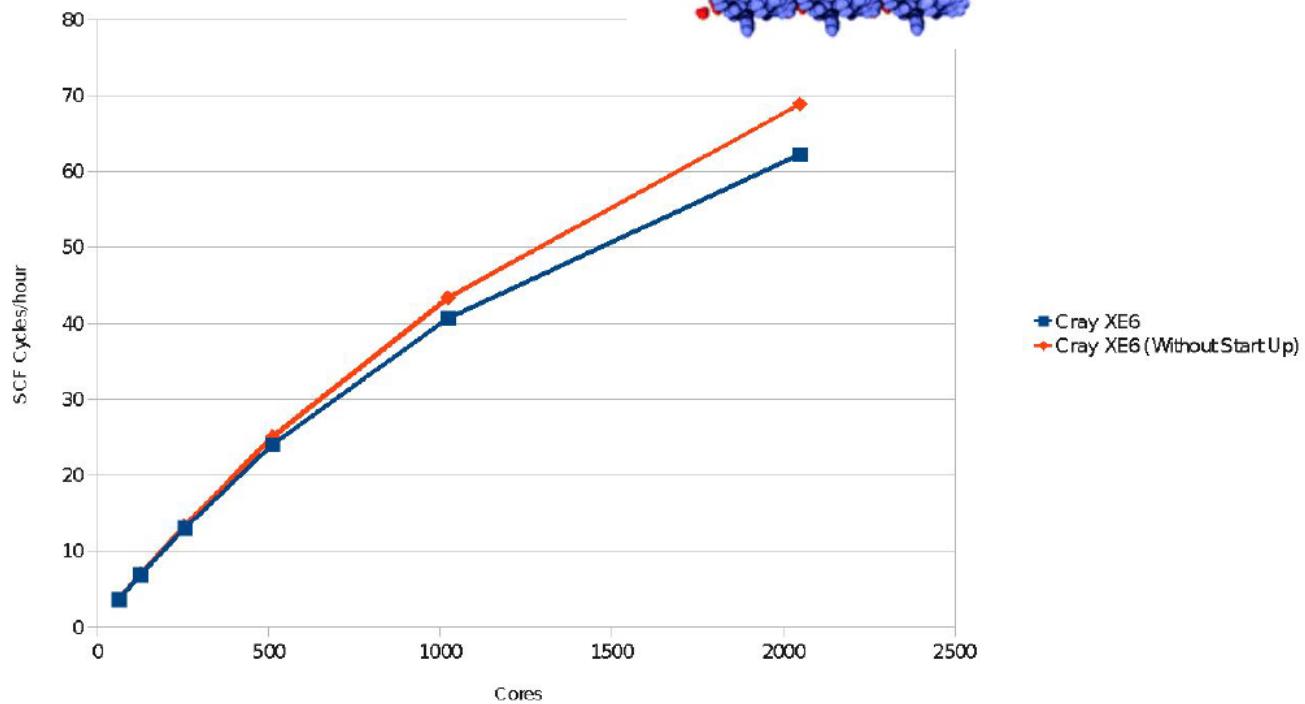
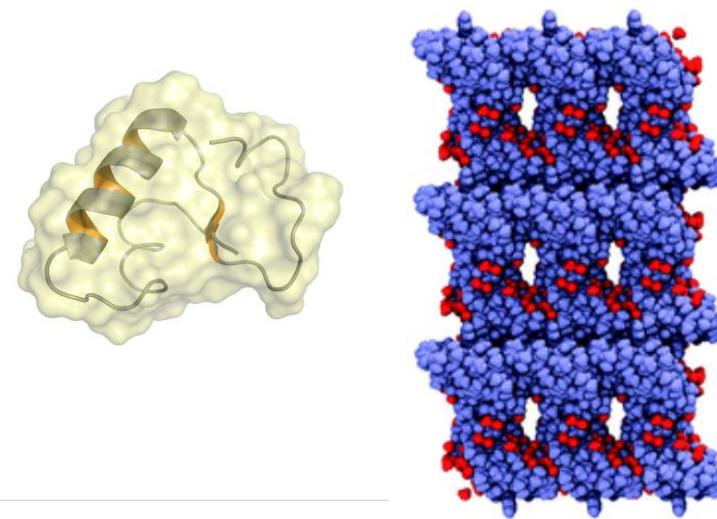


# More Examples

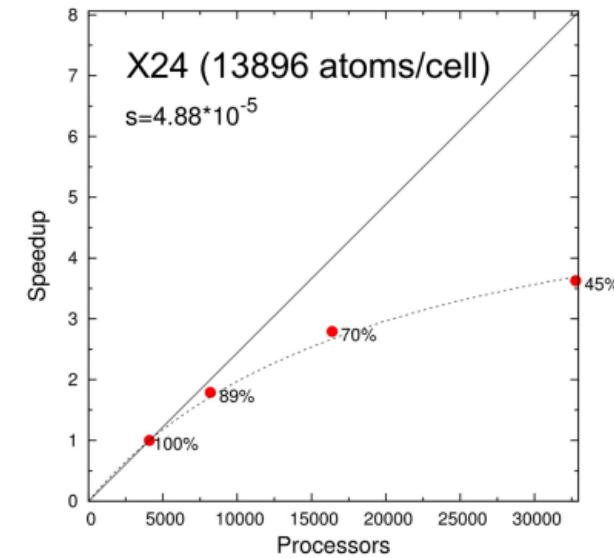
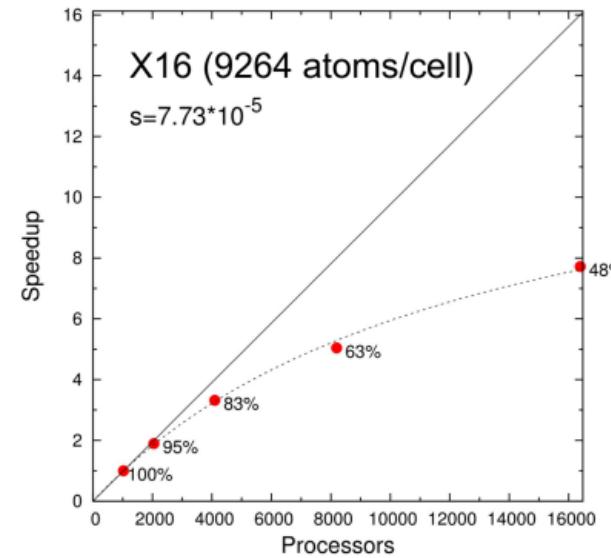
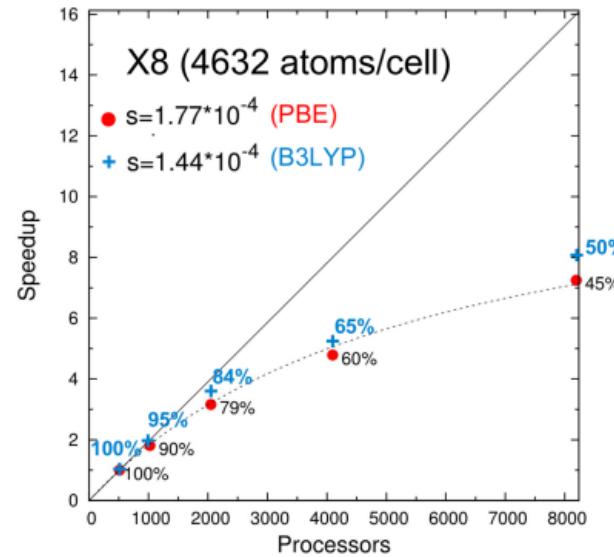
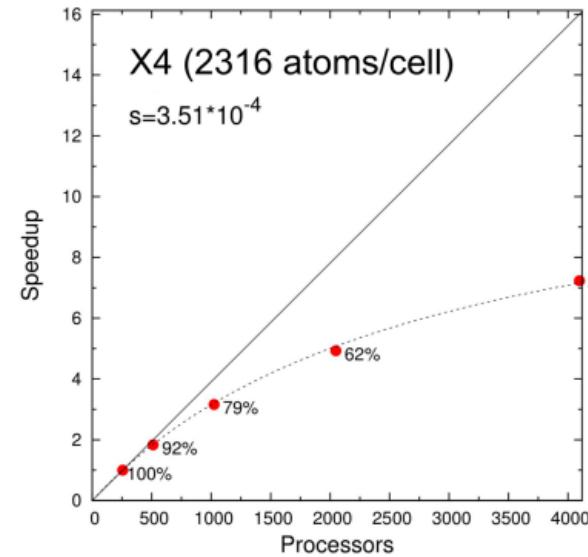
## Crambin

- Small protein (46 residues)
- Crystal structure determined to very high resolution
- **1284 atoms**
- **6-31G\*\* basis set**
- **12354 functions**
- All calculations **B3LYP**
- Latest calculations include a full structural optimisation including water of crystallisation

*Piero Ugliengo et al. Chem. Sci. 2016 7(2)*

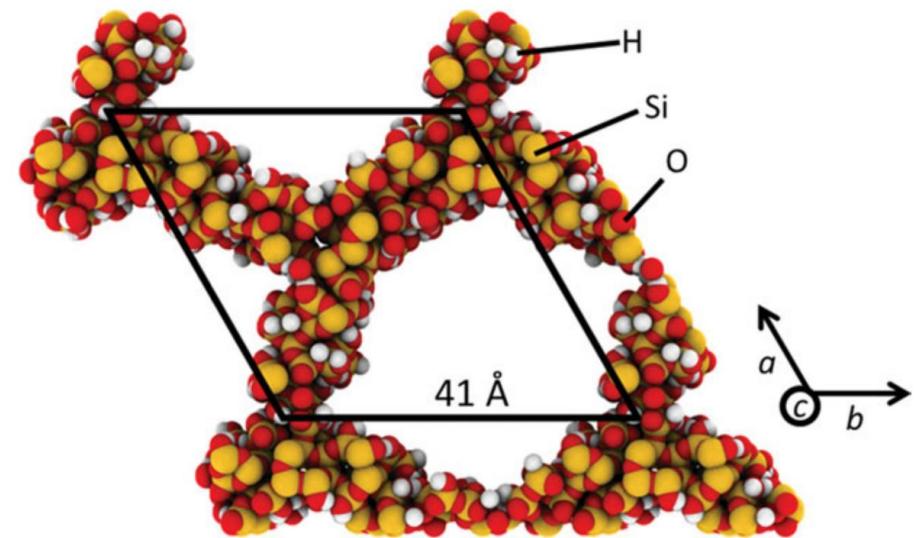


# More Examples



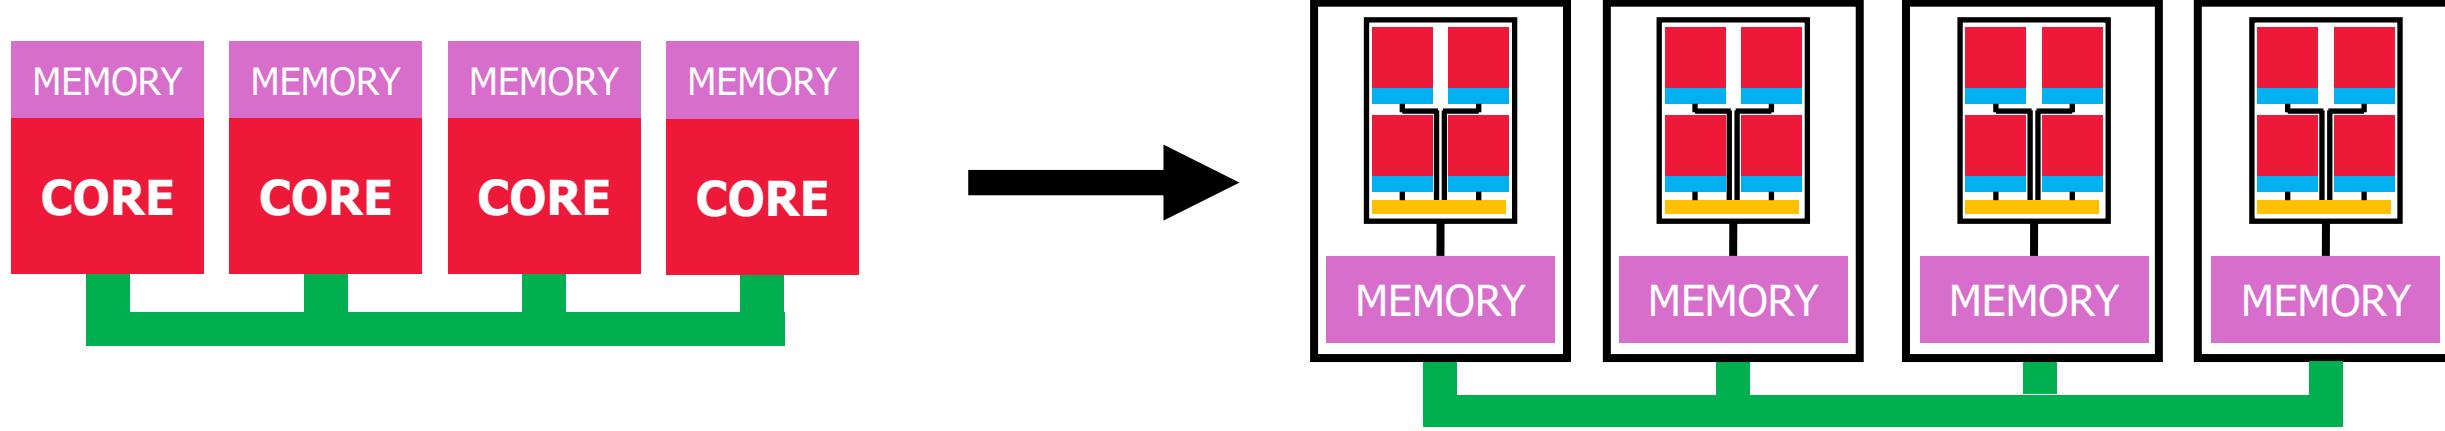
## MCM-41

- **Amorphous (at short range) silica**
- Long range order of pores
- Model developed by Piero Ugliengo et al. at the University of Turin
- **579 atoms, 7790 basis functions**
- Also looked at supercells



# Introduction to OpenMP

Let's go back to the **real hardware** scheme:



**Limits on system size** are posed by **memory/core** available

- PCRYSTAL      hard limit
- MPPCRYSTAL    can use more processes to reduce requirement of each process

The **trick** can be to use **less cores for each node** so that  
**memory/core ratio is increased**

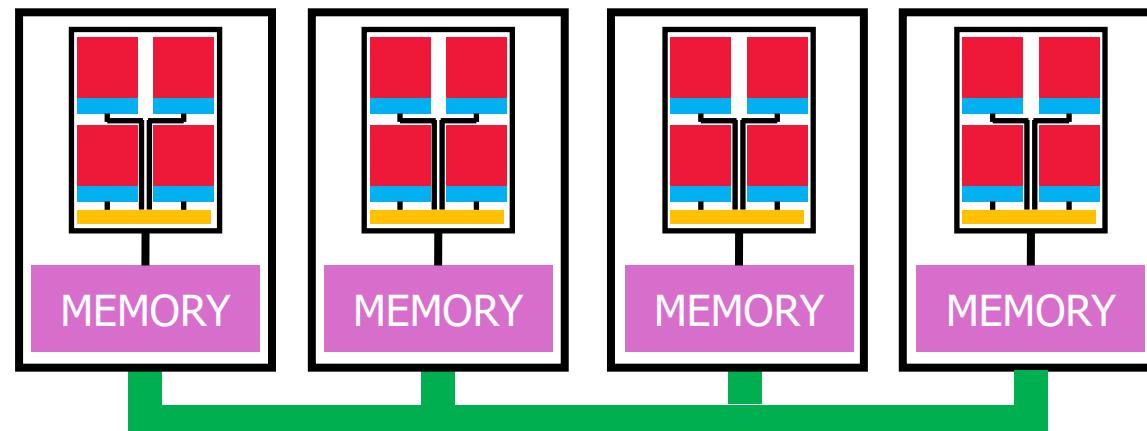
→ But in this way, we are **wasting resources**

# Introduction to OpenMP

The solution is to use **threading**:

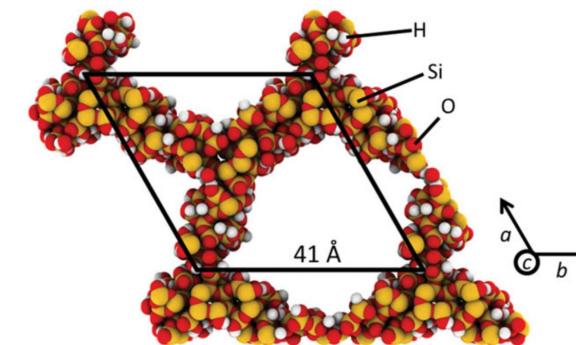
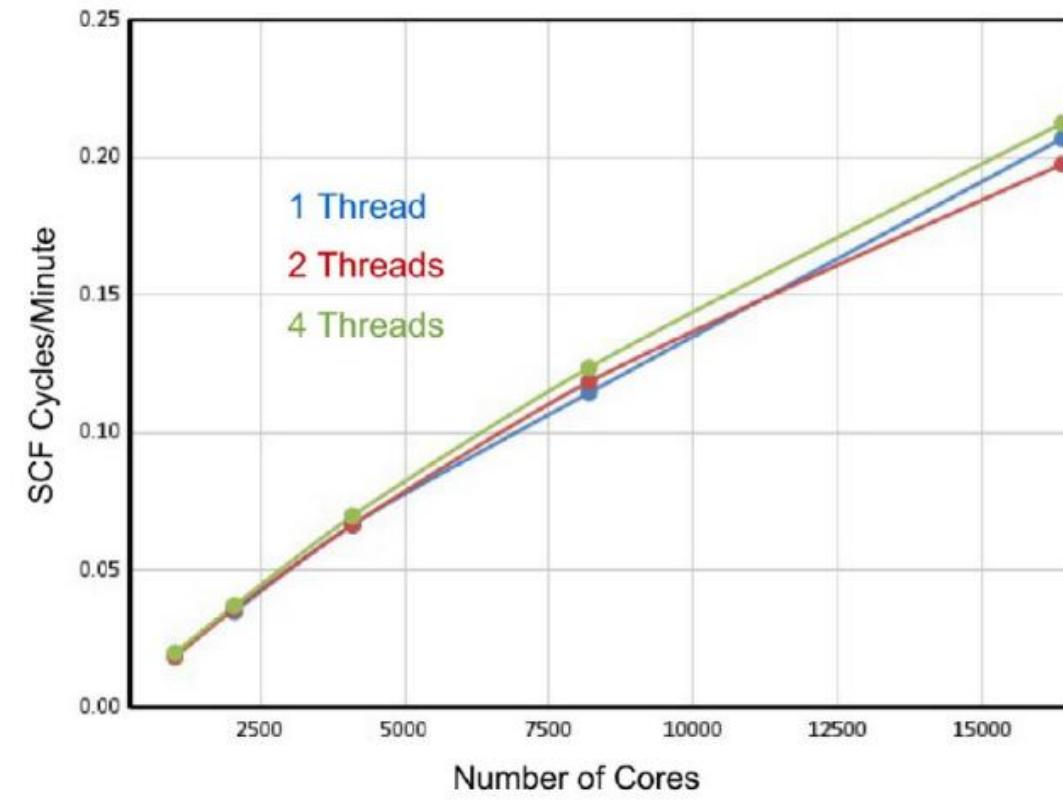
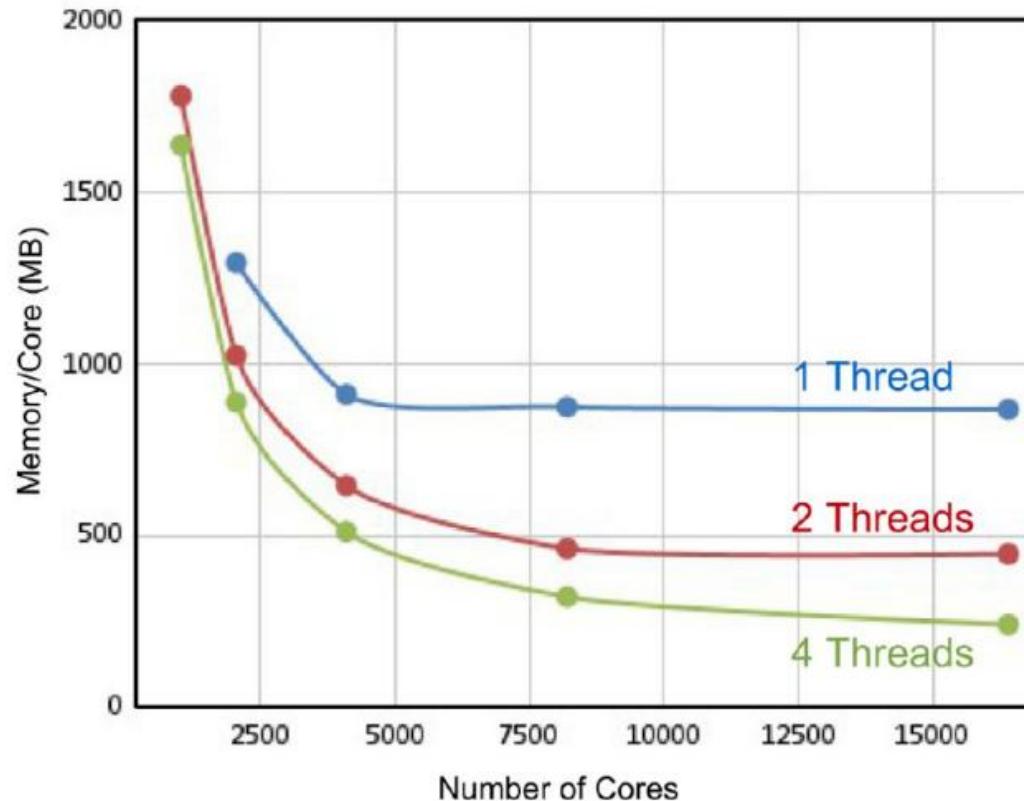
- In specific **regions** of the code, **threads** are generated so to perform **operations in parallel** on **same shared memory**  
→ **No replication of data**
- **This is possible with OpenMP**  
→ New feature in Crystal23
- Can be used only in the same Multicore CPU
- When submitting jobs use:  
`>> export OMP_NUM_THREADS=4`
- Available for **CRYSTAL**, **PCRYSTAL** and **MPPCRYSTAL**

**OpenMP®**



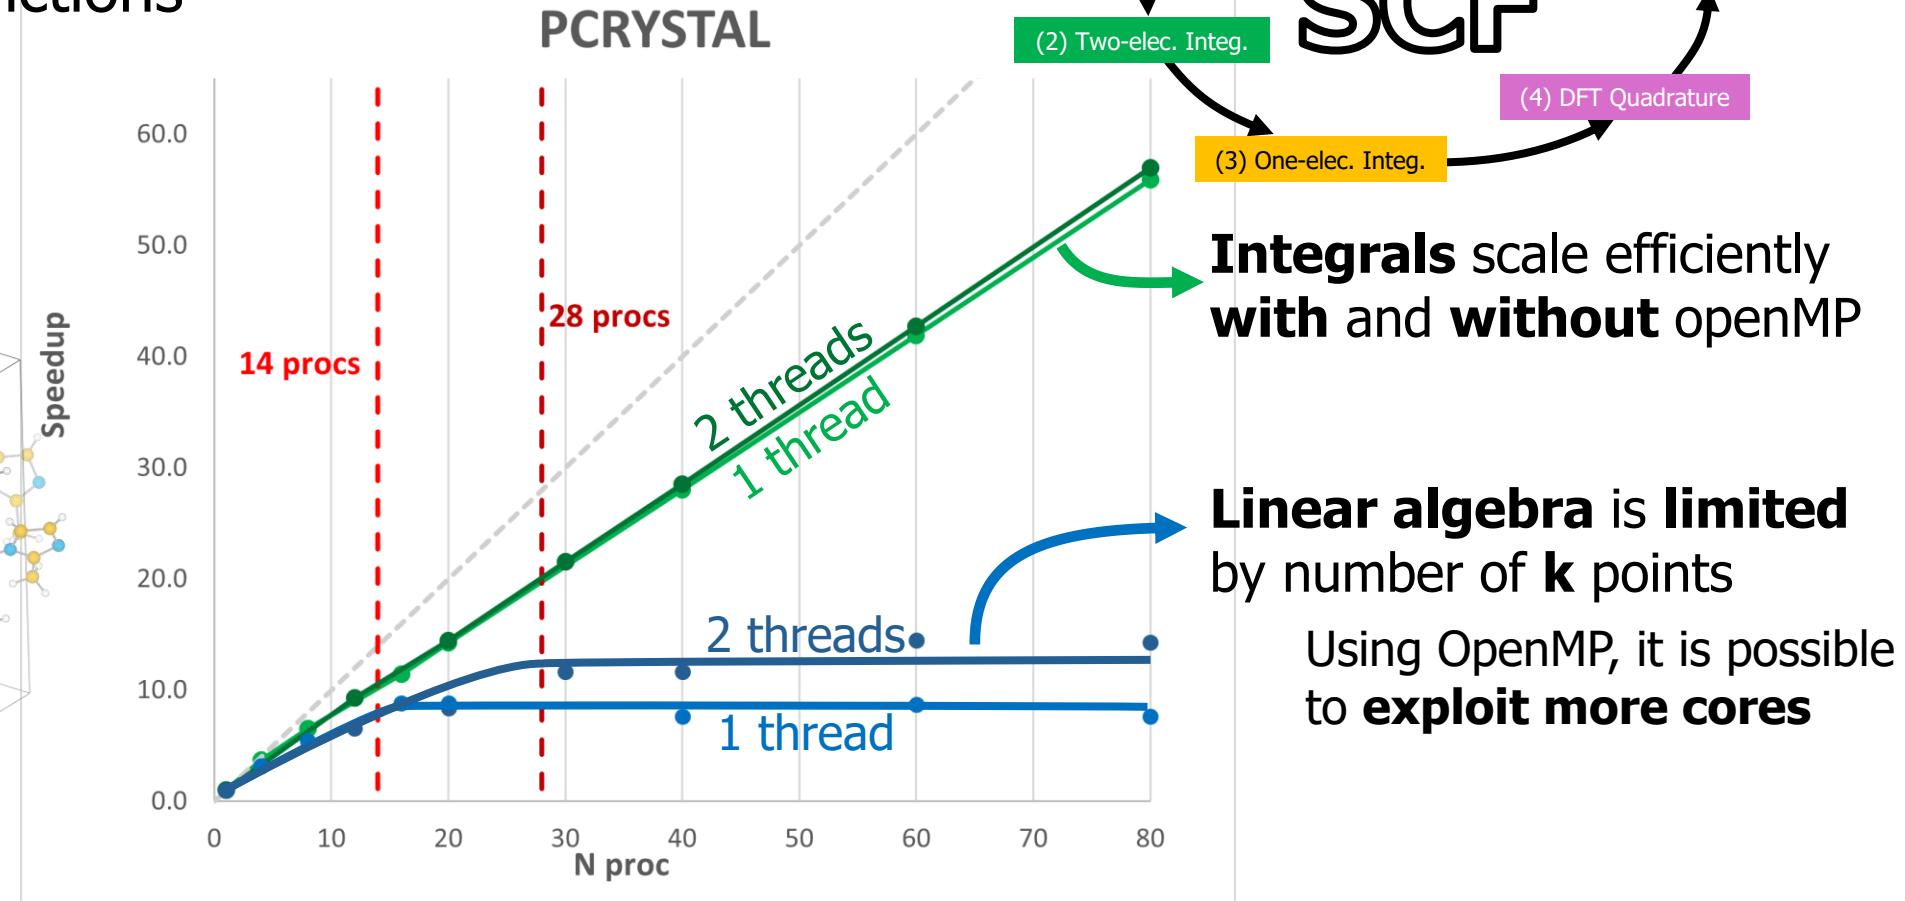
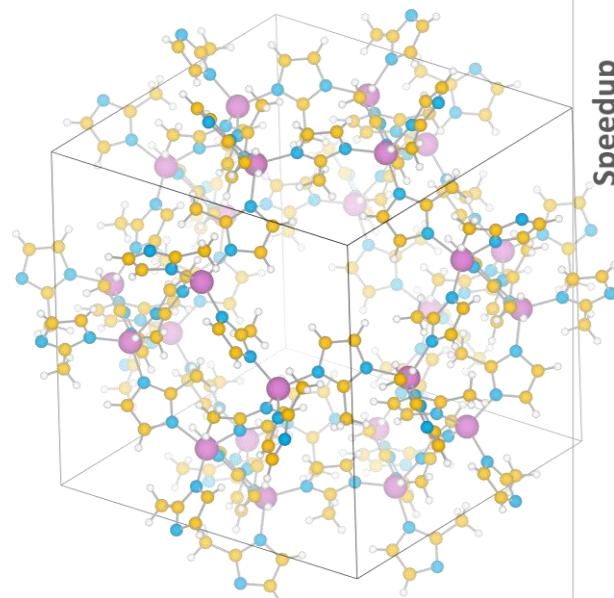
# Example: MCM41x16 - PBE

- MPPCRYSTAL + openMP
- 9264 Atoms
- 124640 Basis Functions



# Example: ZIF-8 with PCRYSTAL

- PCRYSTAL + openMP
- 276 Atoms
- 2772 Basis Functions
- 14  $\mathbf{k}$  points



# GCRYSTAL



- Another option to solve **linear algebra**
- Use **GPUs** to perform **large matrix** diagonalizations
  - full matrix in **GPU memory**
- GPUs are **co-processors**:
  - Used *on top* of CPUs just for some **heavy tasks**
- Gain **access** to **HPC** based on GPU resources
  - Usually required to demonstrate the capability to run on GPUs
- Alternative to solve **very large systems** with **smaller HW setup**
- Ongoing development, probably available in **next release!**

***Just for NVIDIA GPUs for now***

**If want to try alpha version, lets get in touch**

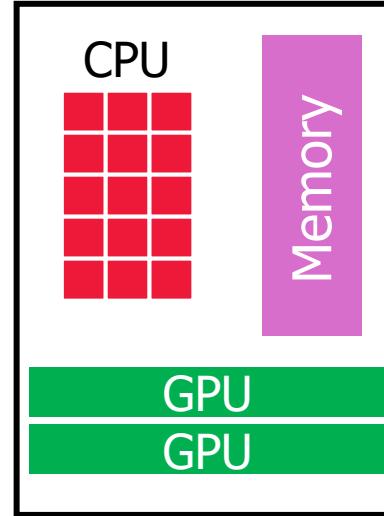
# GCRYSTAL - Logic and Approach

## Common Architectures

- Many CPU cores (**e.g. 128**)
- Few GPUs (**e.g. 4**)

## Hybrid CPU/GPU execution

- Use as many MPI processes as needed **to saturate CPU cores**
- Just bind **few processes** to **GPUs**
- Reorganize **k point distribution**:  
→ only those processes get **k** points

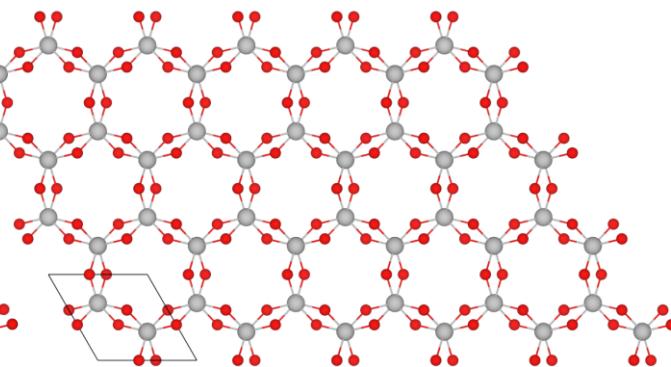
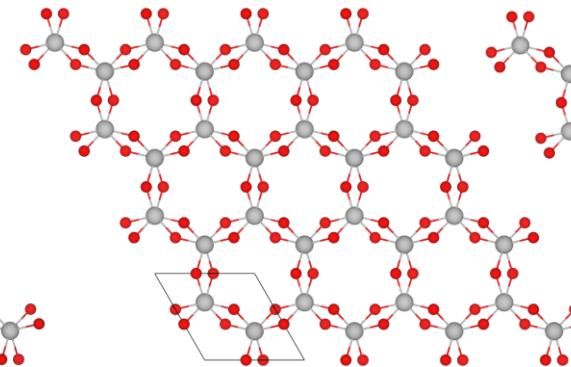
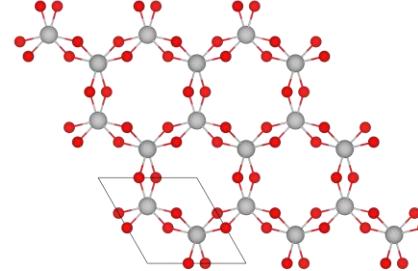
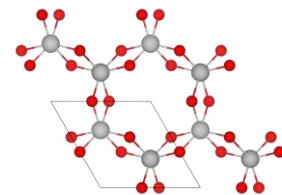


→ It is **more efficient** to perform operations in **series** on **GPU** rather than to **parallelize** on **CPU cores**

# Example $\alpha$ -quartz

## Conditions:

- PBE
- 128 procs
- 0/8 GPUs
- 14  $\mathbf{k}$  points



- 12 Atoms
- 248 BFs

- 96 Atoms
- 1984 BFs

- 324 Atoms
- 6696 BFs

- 648 Atoms
- 13392 BFs

→ LA time in SCF:

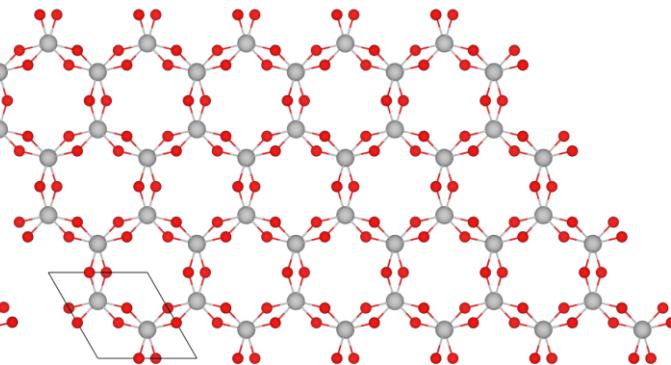
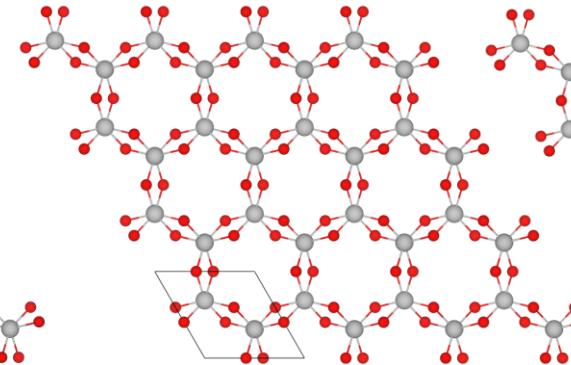
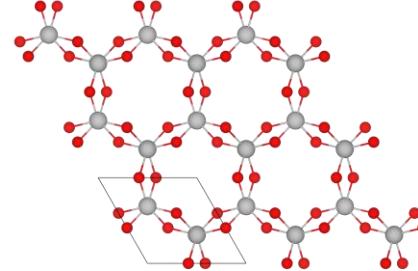
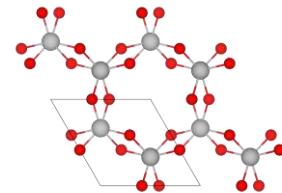
PCRYSTAL	0.07	15.4	606	4454
GCRYSTAL 1 GPU	0.31	2.4 (-84%)	37 (-94%)	228 (-95%)

14 diagonalizations **in series**  
(but more efficient)

# Example $\alpha$ -quartz

## Conditions:

- PBE
- 128 procs
- 0/8 GPUs
- 14  $\mathbf{k}$  points



- 12 Atoms
- 248 BFs

- 96 Atoms
- 1984 BFs

- 324 Atoms
- 6696 BFs

- 648 Atoms
- 13392 BFs

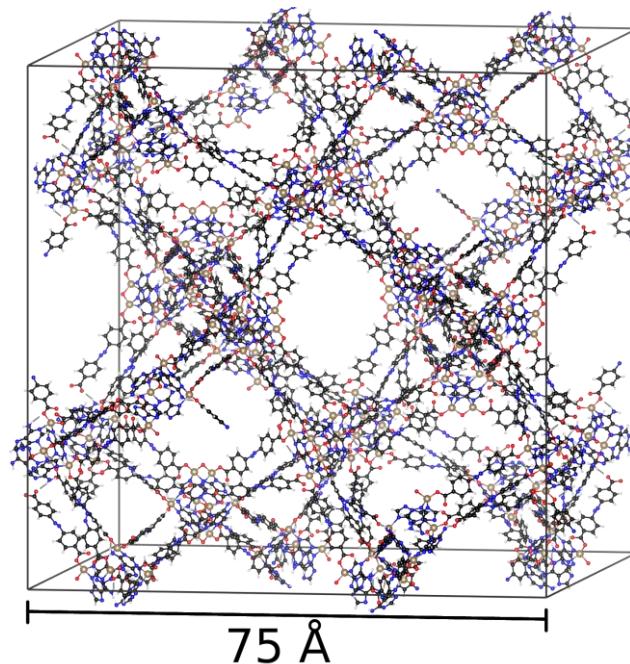
→ LA time in SCF:

PCRYSTAL	0.07	15.4	606	4454
GCRYSTAL 1 GPU	0.31	2.4 (-84%)	37 (-94%)	228 (-95%)
GCRYSTAL 2 GPUs	0.14	1.4 (-91%)	20 (-97%)	116 (-97%)
GCRYSTAL 4 GPUs	0.12	1.0 (-93%)	14 (-98%)	74 (-98%)
GCRYSTAL 8 GPUs	0.11	0.8 (-94%)	10 (-98%)	51 (-99%)

# Example bio-MOF-102

## Conditions:

- 2808 Atoms/Cell
- 33672 Basis Functions
- PBEsol0
- 1 k Point
- 0-1 GPU



→ LA time in SCF:

MPPCRYSTAL	1 Node (128 cores)	538
MPPCRYSTAL	2 Nodes (256 cores)	288
MPPCRYSTAL	4 Nodes (512 cores)	147
MPPCRYSTAL	8 Nodes (1024 cores)	60
GCRYSTAL	1 Node (128 cores + 1 GPU)	87

G. Ambrogio, et al. "Accelerated linear algebra for large scale DFT calculations of materials on CPU/GPU architectures with CRYSTAL" JCP 162.8 (2025)

# Final Remarks

2 parallelisation strategies currently available:

- **PCRYSTAL**

- Good for **medium size systems**
- **Scalability limited** by **k** points
- Can't study systems bigger than serial CRYSTAL

- **MPPCRYSTAL**

- Good for **large systems**
- Can **scale well** for very big problems

**OpenMP** can be used with both

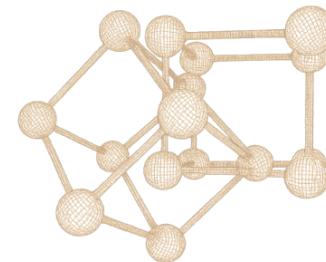
- Reduces memory required
- Improve efficiency just for PCRYSTAL, if limited by **k** points

In all cases, a little bit of attention to the **input** can **improve the performance**

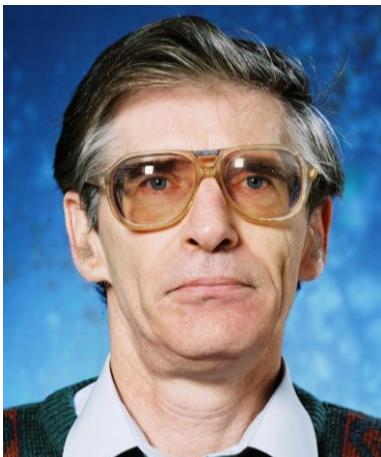
- **GCRYSTAL** (*coming soon*)

- Use **GPUs** for Linear Algebra
- Hopefully will be extended to other sections of the code

# Acknowledgment



Advanced School on  
**QUANTUM MODELLING**  
of Materials with CRYSTAL



**Vic Saunders**



**Dr. Ian J. Bush**

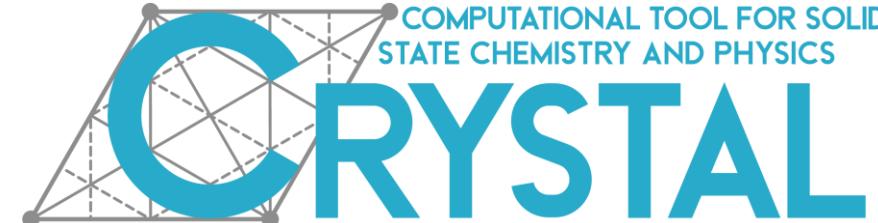


**Dr. Barry Searle**

Part of the material presented has been adapted  
from slides kindly provided by Dr. **Ian Bush**



UK Research  
and Innovation



# SCF in CRYSTAL

P  
MPP

Advanced School on  
**QUANTUM MODELLING**  
of Materials with CRYSTAL



(2) Two-elec. Integ.

$$E_{\text{coul}} = \frac{1}{2} \sum_{\mu\nu} \sum_{\mathbf{g}} P_{\mu\nu}^{\mathbf{g}} \sum_{\sigma\rho} \sum_{\mathbf{q}} P_{\sigma\rho}^{\mathbf{q}} \sum_{\mathbf{s}} (\chi_{\mu}^{\mathbf{0}} \chi_{\nu}^{\mathbf{g}} | \chi_{\sigma}^{\mathbf{s}} \chi_{\rho}^{\mathbf{s+q}})$$

$$E_{\text{exx}} = -\frac{1}{2} \sum_{\mu\nu} \sum_{\mathbf{g}} P_{\mu\nu}^{\mathbf{g}} \sum_{\sigma\rho} \sum_{\mathbf{q}} P_{\sigma\rho}^{\mathbf{q}} \sum_{\mathbf{s}} (\chi_{\mu}^{\mathbf{0}} \chi_{\sigma}^{\mathbf{s}} | \chi_{\nu}^{\mathbf{g}} \chi_{\rho}^{\mathbf{s+q}})$$

(3) One-elec. Integ.

$$E_{\text{kin}} = \sum_{\mu\nu} \sum_{\mathbf{g}} P_{\mu\nu}^{\mathbf{g}} \langle \chi_{\mu}^{\mathbf{0}} | -\frac{\nabla^2}{2} | \chi_{\nu}^{\mathbf{g}} \rangle$$

$$E_{\text{e-nucl}} = - \sum_{\mu\nu} \sum_{\mathbf{g}} P_{\mu\nu}^{\mathbf{g}} \sum_{\lambda} Z_{\lambda} \sum_{\mathbf{h}} \left\langle \chi_{\mu}^{\mathbf{0}} \right| \frac{1}{|\mathbf{r} - \mathbf{h} - \mathbf{s}_Z|} \left| \chi_{\nu}^{\mathbf{g}} \right\rangle$$

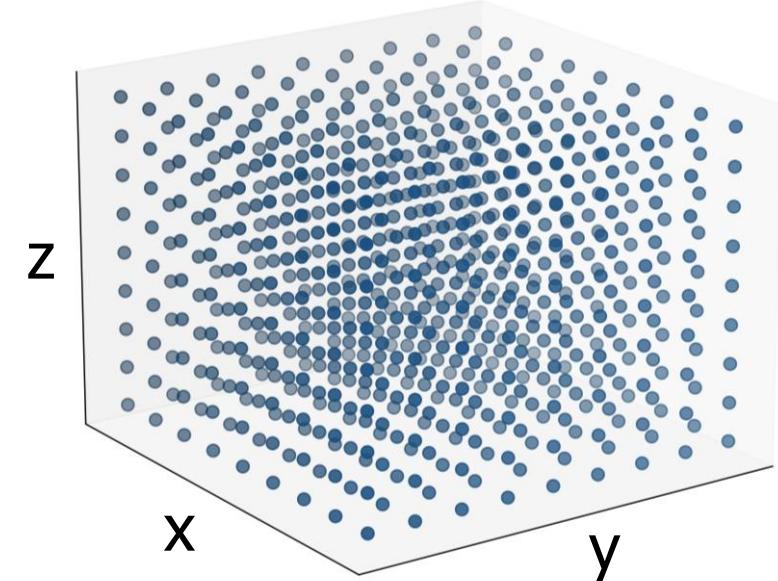
# SCF in CRYSTAL

## (4) DFT Quadrature

$$E_{xc}[\rho(\mathbf{r})] = \int d\mathbf{r} \varepsilon_c[\rho(\mathbf{r})] + \int d\mathbf{r} \varepsilon_x[\rho(\mathbf{r})]$$

Batches of Points  
Distributed to  
Processes

Each Point is Independent  
from Others



Computed on a 3D  
Grid of Points



# SCF in CRYSTAL



P

MPP



UNIVERSITÀ  
DI TORINO

(5) Linear Algebra

$$H_{\mu\nu}^{\mathbf{k}} = \sum_{\mathbf{g}} H_{\mu\nu}^{\mathbf{g}} e^{i\mathbf{k}\cdot\mathbf{g}}$$

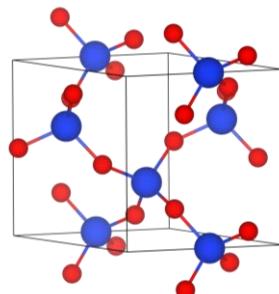
For each  $\mathbf{k}$ :

$$\mathbf{H}^{\mathbf{k}} \mathbf{A}^{\mathbf{k}} = \mathbf{S}^{\mathbf{k}} \mathbf{A}^{\mathbf{k}} \mathbf{E}^{\mathbf{k}}$$

n AOs



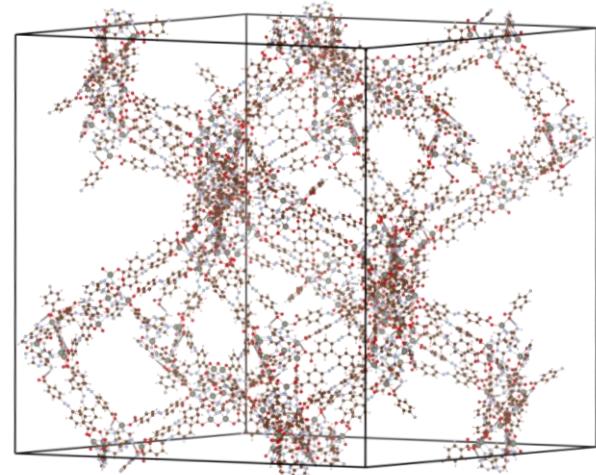
Small Unit Cell



- Few Basis Functions
- Many  $\mathbf{k}$  Points

**Efficient Parallelization**

Large Unit Cell



- Many Basis Functions
- Few  $\mathbf{k}$  Points

**Poor Parallelization**

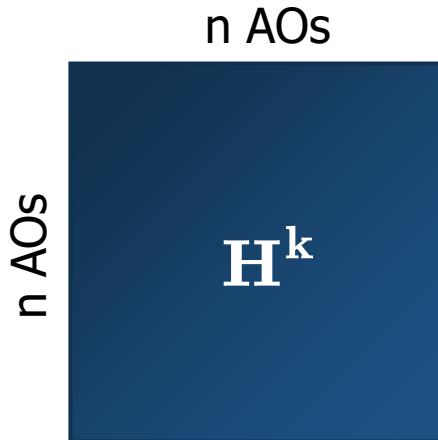
# SCF in CRYSTAL

(5) Linear Algebra

$$H_{\mu\nu}^{\mathbf{k}} = \sum_{\mathbf{g}} H_{\mu\nu}^{\mathbf{g}} e^{i\mathbf{k}\cdot\mathbf{g}}$$

For each  $\mathbf{k}$ :

$$\mathbf{H}^{\mathbf{k}} \mathbf{A}^{\mathbf{k}} = \mathbf{S}^{\mathbf{k}} \mathbf{A}^{\mathbf{k}} \mathbf{E}^{\mathbf{k}}$$



P  
MPP



Advanced School on  
**QUANTUM MODELLING**  
of Materials with CRYSTAL



UNIVERSITÀ  
DI TORINO

**General Rule**

$n \mathbf{k}$  points  $\geq n$  processes

**Large Systems**

**Large-Scale Architectures**

# SCF in CRYSTAL



P

MPP



UNIVERSITÀ  
DI TORINO

(5) Linear Algebra

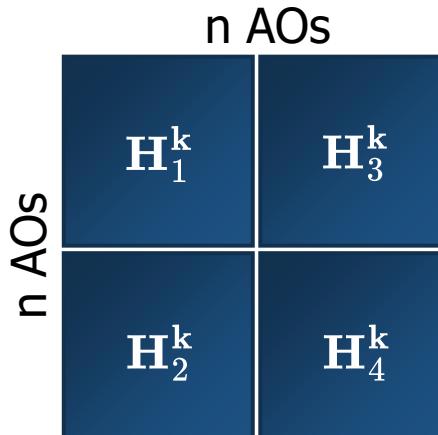
General Rule

$$H_{\mu\nu}^{\mathbf{k}} = \sum_{\mathbf{g}} H_{\mu\nu}^{\mathbf{g}} e^{i\mathbf{k}\cdot\mathbf{g}}$$

$n$   $\mathbf{k}$  points  $\geq n$  processes

For each  $\mathbf{k}$ :

$$\mathbf{H}^{\mathbf{k}} \mathbf{A}^{\mathbf{k}} = \mathbf{S}^{\mathbf{k}} \mathbf{A}^{\mathbf{k}} \mathbf{E}^{\mathbf{k}}$$



Large Systems

Large-Scale Architectures

**MPPCRYSTAL** → **ScaLAPACK**

- Allocation of Large Matrices Across Multiple Processes
- Linear Algebra Operations Executed by Multiple Processes (using MPI)

R. Orlando, M. Delle Piane, I. J. Bush, P. Ugliengo, M. Ferrabone, and R. Dovesi, A new massively parallel version of CRYSTAL for large systems on high performance computing architectures, *J. Comput. Chem.*, **2012**, 33.28, 2276-2284

# SCF in CRYSTAL

(6) Fermi Surface

$$n_{\text{elect.}} = \frac{1}{\Omega} \int_{\text{FBZ}} d\mathbf{k} \sum_i \theta[\epsilon_F - \varepsilon_i^{\mathbf{k}}]$$

Not parallelized

(1) Density Matrix

$$P_{\sigma\rho}^{\mathbf{g}} = \frac{1}{\Omega} \sum_i f_i \sum_{\mathbf{k}} e^{i\mathbf{k}\cdot\mathbf{g}} a_{\sigma,i}^{\mathbf{k}} [a_{\rho,i}^{\mathbf{k}}]^* \theta[\epsilon_F - \varepsilon_i^{\mathbf{k}}]$$

Parallelized over  $\mathbf{k}$  points

**Not expensive**